


More frames don't help: overall flat then falling — temporal-ordering categories degrade most

2×2-stitch presentation · InternVL3-8B · 1000 CVBench questions × 4 passes · line = mean, band = ±1 std, dot = best budget · panels sorted by gain 8 → 128

Dashed = chance, the score of a uniform random guesser: mean over that panel's questions of $1/(\text{number of answer options})$. Not a flat 25% — Cross-video Scene Recognition is mostly yes/no.

accuracy (%)



frame budget (stitched 2×2 timesteps per question, log scale)